

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: df = pd.read_csv('data/credit_card_default.csv')
df.head(10)
```

Out []:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_1	PAY_2	PAY_3	PAY_4	...	BILL_AMT4	BILL_AMT5	BILL_AMT6
0	1	20000	2	2	1	24	2	2	-1	-1	...	0	0	0
1	2	120000	2	2	2	26	-1	2	0	0	...	3272	3455	3261
2	3	90000	2	2	2	34	0	0	0	0	...	14331	14948	15549
3	4	50000	2	2	1	37	0	0	0	0	...	28314	28959	29547
4	5	50000	1	2	1	57	-1	0	-1	0	...	20940	19146	19131
5	6	50000	1	1	2	37	0	0	0	0	...	19394	19619	20024
6	7	500000	1	1	2	29	0	0	0	0	...	542653	483003	473944
7	8	100000	2	2	2	23	0	-1	-1	0	...	221	-159	567
8	9	140000	2	3	1	28	0	0	2	0	...	12211	11793	3719
9	10	20000	1	3	2	35	-2	-2	-2	-2	...	0	13007	13912

10 rows × 25 columns

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    30000 non-null  int64
1   LIMIT_BAL                            30000 non-null  int64
2   SEX                                  30000 non-null  int64
3   EDUCATION                            30000 non-null  int64
4   MARRIAGE                             30000 non-null  int64
5   AGE                                  30000 non-null  int64
6   PAY_1                                30000 non-null  int64
7   PAY_2                                30000 non-null  int64
8   PAY_3                                30000 non-null  int64
9   PAY_4                                30000 non-null  int64
10  PAY_5                                30000 non-null  int64
11  PAY_6                                30000 non-null  int64
12  BILL_AMT1                            30000 non-null  int64
13  BILL_AMT2                            30000 non-null  int64
14  BILL_AMT3                            30000 non-null  int64
15  BILL_AMT4                            30000 non-null  int64
16  BILL_AMT5                            30000 non-null  int64
17  BILL_AMT6                            30000 non-null  int64
18  PAY_AMT1                             30000 non-null  int64
19  PAY_AMT2                             30000 non-null  int64
20  PAY_AMT3                             30000 non-null  int64
21  PAY_AMT4                             30000 non-null  int64
22  PAY_AMT5                             30000 non-null  int64
23  PAY_AMT6                             30000 non-null  int64
24  default payment next month           30000 non-null  int64
dtypes: int64(25)
memory usage: 5.7 MB
```

```
In [ ]: df.rename(columns = lambda x: x.lower(), inplace=True)
df.rename(columns = {"default payment next month": "default"}, inplace=True)
df.columns
```

```
Out [ ]: Index(['id', 'limit_bal', 'sex', 'education', 'marriage', 'age', 'pay_1',
            'pay_2', 'pay_3', 'pay_4', 'pay_5', 'pay_6', 'bill_amt1', 'bill_amt2',
            'bill_amt3', 'bill_amt4', 'bill_amt5', 'bill_amt6', 'pay_amt1',
            'pay_amt2', 'pay_amt3', 'pay_amt4', 'pay_amt5', 'pay_amt6', 'default'],
            dtype='object')
```

```
In [ ]: df.nunique()
```

```
Out [ ]: id          30000
limit_bal      81
sex            2
education      7
marriage       3
age           56
pay_1         11
pay_2         11
pay_3         11
pay_4         11
pay_5         10
pay_6         10
bill_amt1     22723
bill_amt2     22346
bill_amt3     22026
bill_amt4     21548
bill_amt5     21010
bill_amt6     20604
pay_amt1      7943
pay_amt2      7899
pay_amt3      7518
pay_amt4      6937
pay_amt5      6897
pay_amt6      6939
default       2
dtype: int64
```

```
In [ ]: df["male"] = (df['sex'] == 1).astype(int)
df['grad_school'] = (df['education']==1).astype(int)
df["university"] = (df['education']==2).astype(int)
df["married"] = (df['marriage']==1).astype(int)
```

```
In [ ]: bill_amt_features = ["bill_amt" + str(i) for i in range(1,7)]
pay_amt_features = ["pay_amt" + str(i) for i in range(1,7)]
```

```
In [ ]: binary_features = ['male', 'grad_school', 'university', 'married']
pay_features = ["pay_" + str(i) for i in range(1,7)]
num_features = ["limit_bal", "age"] + bill_amt_features + pay_amt_features + pay_features
```

```
In [ ]: x = df[num_features + binary_features]
y = df['default']
```

```
In [ ]: x
```

```
Out [ ]:
```

	limit_bal	age	bill_amt1	bill_amt2	bill_amt3	bill_amt4	bill_amt5	bill_amt6	pay_amt1	pay_amt2	...	pay_1	pay_2	pay_3
0	20000	24	3913	3102	689	0	0	0	0	689	...	2	2	2
1	120000	26	2682	1725	2682	3272	3455	3261	0	1000	...	-1	2	2
2	90000	34	29239	14027	13559	14331	14948	15549	1518	1500	...	0	0	0
3	50000	37	46990	48233	49291	28314	28959	29547	2000	2019	...	0	0	0
4	50000	57	8617	5670	35835	20940	19146	19131	2000	36681	...	-1	0	0
...
29995	220000	39	188948	192815	208365	88004	31237	15980	8500	20000	...	0	0	0
29996	150000	43	1683	1828	3502	8979	5190	0	1837	3526	...	-1	-1	-1
29997	30000	37	3565	3356	2758	20878	20582	19357	0	0	...	4	3	3
29998	80000	41	-1645	78379	76304	52774	11855	48944	85900	3409	...	1	-1	-1
29999	50000	46	47929	48905	49764	36535	32428	15313	2078	1800	...	0	0	0

30000 rows × 24 columns

```
In [ ]: from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, RepeatedStratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score
```

```
In [ ]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state = 134)
```

```
In [ ]: sc_features = ["limit_bal", "age"] + bill_amt_features + pay_amt_features
```

```
In [ ]: sc = StandardScaler()
```

```
In [ ]: x_train.loc[:,sc_features] = sc.fit_transform(x_train[sc_features])
x_test.loc[:,sc_features] = sc.transform(x_test[sc_features])
```

```
In [ ]: x_train.head()
```

```
Out [ ]:
```

	limit_bal	age	bill_amt1	bill_amt2	bill_amt3	bill_amt4	bill_amt5	bill_amt6	pay_amt1	pay_amt2	...	pay_1	pay
29245	-0.911109	-1.136689	-0.055124	-0.613938	-0.591388	-0.569915	-0.523784	-0.584140	-0.221162	-0.183388	...	0	
26498	-0.911109	1.790000	-0.039584	-0.001461	-0.127831	-0.672645	-0.490840	-0.652265	-0.219044	-0.177226	...	0	
20130	-0.678940	0.814437	0.238241	0.293119	0.358289	0.454690	0.548120	0.609451	-0.154595	-0.097121	...	2	
25126	4.428778	1.573208	3.676412	3.754222	3.891379	3.946858	4.047345	3.870115	0.444513	0.297247	...	0	
20337	-0.911109	2.765562	0.021160	0.029323	0.068152	0.112776	-0.165604	-0.143091	-0.341952	-0.148635	...	3	

5 rows × 24 columns

KNN Classifier

```
In [ ]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [ ]: # commented for being computationally intensive
'''param_grid_knn = {
    'n_neighbors': np.arange(5,15),
    'p': (1,2),
    'weights': ('uniform', 'distance'),
    'metric': ('minkowski', 'chebyshev')
}
knn_gscv = KNeighborsClassifier()
knn_gscv = GridSearchCV(knn_gscv, param_grid=param_grid_knn, scoring='accuracy', cv=5)
knn_gscv.fit(x_train, y_train)'''
```

```
Out [ ]: "param_grid_knn = {\n    'n_neighbors': np.arange(5,15),\n    'p': (1,2),\n    'weights': ('uniform', 'distance'),\n    'metric': ('minkowski', 'chebyshev')\n}\n\nknn_gscv = KNeighborsClassifier()\nknn_gscv = GridSearchCV(knn_gscv, param_grid=param_grid_knn, scoring='accuracy', cv=5)\nknn_gscv.fit(x_train, y_train)"
```

```
In [ ]: #print("Best hyperparameters:", knn_gscv.best_params_)
#print("Best accuracy score:", knn_gscv.best_score_)
```

Best hyperparameters: {'metric': 'minkowski', 'n_neighbors': 12, 'p': 2, 'weights': 'uniform'}

Best accuracy score: 0.8124583333333334

```
In [ ]: knn = KNeighborsClassifier(n_neighbors=12, weights='uniform', p=2, metric='minkowski')
knn.fit(x_train, y_train)
```

```
Out [ ]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=12)
```

```
In [ ]: y_pred_knn = knn.predict(x_test)
```

```
In [ ]: print("Accuracy : " + str(accuracy_score(y_test, y_pred_knn)))
```

Accuracy : 0.8105

Support Vector Classification

```
In [ ]: from sklearn.svm import SVC
```

```
In [ ]: # commented because it took a whopping 84 minutes to execute!
'''
param_grid_svm = {'C': [0.1, 1, 10, 100, 1000],
                  'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                  'kernel': ['rbf']}

svm_gscv = GridSearchCV(SVC(), param_grid_svm, refit = True, verbose = 3)

# fitting the model for grid search
svm_gscv.fit(x_train, y_train)
'''
```

```

Out [ ]: "\nparam_grid_svm = {'C': [0.1, 1, 10, 100, 1000],\n\t\t\t\t\t'gamma': [1, 0.1, 0.01, 0.001, 0.0001],\n\t\t\t\t\t'kernel': ['rbf']}\n\nsvm_gscv = GridSearchCV(SVC(), param_grid_svm, refit = True, verbose = 3)\n\n# fitting the model for grid search\nsvm_gscv.fit(x_train, y_train)\n"

In [ ]: #print("Best hyperparameters:", svm_gscv.best_params_)
        #print("Best accuracy score:", svm_gscv.best_score_)

Best hyperparameters: {'C': 1000, 'gamma': 0.001, 'kernel': 'rbf'}

Best accuracy score: 0.8206249999999999

In [ ]: sv = SVC(C=1000, gamma= 0.001, kernel='rbf')
        sv.fit(x_train, y_train)

Out [ ]: ▼ SVC
        SVC(C=1000, gamma=0.001)

In [ ]: y_pred_svc = sv.predict(x_test)

In [ ]: print("Accuracy : " + str(accuracy_score(y_test, y_pred_svc)))

Accuracy : 0.817

```

Logistic Regression Classifier

```

In [ ]: from sklearn.linear_model import LogisticRegression

In [ ]: param_grid_log = { 'solver': ['newton-cg', 'lbfgs', 'liblinear'],
                          'penalty': ['l2'],
                          'C' : [100, 10, 1.0, 0.1, 0.01] }

cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
log_gscv = GridSearchCV(LogisticRegression(), param_grid_log, n_jobs=-1, cv=cv, scoring='accuracy', error_score=0)
log_gscv.fit(x_train, y_train)

Out [ ]: ▶ GridSearchCV
        ▶ estimator: LogisticRegression
          ▶ LogisticRegression

In [ ]: print("Best hyperparameters:", log_gscv.best_params_)
        print("Best accuracy score:", log_gscv.best_score_)

Best hyperparameters: {'C': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
Best accuracy score: 0.8120972222222222

In [ ]: log = LogisticRegression(C=100, penalty='l2', solver='lbfgs')
        log.fit(x_train, y_train)

Out [ ]: ▼ LogisticRegression
        LogisticRegression(C=100)

In [ ]: y_pred_log = log.predict(x_test)

In [ ]: print("Accuracy : " + str(accuracy_score(y_test, y_pred_log)))

Accuracy : 0.8053333333333333

```

Bagging Classifier

```

In [ ]: from sklearn.ensemble import BaggingClassifier

In [ ]: '''param_grid_bag = { 'n_estimators': [10, 100, 1000] }

#cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
bag_gscv = GridSearchCV(BaggingClassifier(), param_grid_bag, n_jobs=-1, cv=5, scoring='accuracy', error_score=0)
bag_gscv.fit(x_train, y_train)'''

```

```
Out [ ]: ▶ GridSearchCV  
▶ estimator: BaggingClassifier  
▶ BaggingClassifier
```

```
In [ ]: #print("Best hyperparameters:", bag_gscv.best_params_)  
#print("Best accuracy score:", bag_gscv.best_score_)
```

Best hyperparameters: {'n_estimators': 1000}

Best accuracy score: 0.8165416666666667

Best hyperparameters: {'n_estimators': 1000} Best accuracy score: 0.8165416666666667

```
In [ ]: bag = BaggingClassifier(n_estimators=1000)  
bag.fit(x_train, y_train)
```

```
Out [ ]: ▼ BaggingClassifier  
BaggingClassifier(n_estimators=1000)
```

```
In [ ]: y_pred_bag = bag.predict(x_test)
```

```
In [ ]: print("Accuracy : " + str(accuracy_score(y_test, y_pred_bag)))
```

Accuracy : 0.8143333333333334